

# MongoDB

Or how I learned to stop worrying and love the database

Mathias Stearn

10gen

Bay Area Hadoop Meetup – February 17, 2010

*The resulting [MongoDB] application has literally changed the way the pharma company conducts business. Whereas in the past, patient queries could take minutes to hours, results are now essentially real-time.*

[http://news.cnet.com/8301-13846\\_3-10451248-62.html](http://news.cnet.com/8301-13846_3-10451248-62.html)

*Why bother [with] memcached for caching HTTP sessions when you have an authoritative MongoDB? The performance is there.*

Carl Byström, @cgbystrom on twitter

*Compared to hadoop, Mongo's speed and startup time make developing new queries much easier; what took us two weeks to get working on hadoop was done in two days on mongo.*

Emmett Shear, CTO (and developer) at Justin.tv

*It took me half a day to go from not touching MongoDB to writing some fairly good functionality against it. It makes setting up, configuring, and interfacing with MySQL look archaic – ridiculously archaic.*

[http://geekaustin.org/2010/01/31/  
mongodb-day-geek-austin-data-series](http://geekaustin.org/2010/01/31/mongodb-day-geek-austin-data-series)

- 1 What is MongoDB?
  - Document Oriented
  - JavaScript Enabled
  - Fast, Scalable, Available, and Reliable
- 2 What Makes Mongo Special?
  - Native Language Integration
  - Rich Data Types
  - Atomic Modifiers
  - Dynamic Queries
- 3 MapReduce
  - Built-In MapReduce
  - Easy Hadoop-Mongo Integration
  - Better Hadoop-Mongo Integration

- 1 What is MongoDB?
  - Document Oriented
  - JavaScript Enabled
  - Fast, Scalable, Available, and Reliable
- 2 What Makes Mongo Special?
  - Native Language Integration
  - Rich Data Types
  - Atomic Modifiers
  - Dynamic Queries
- 3 MapReduce
  - Built-In MapReduce
  - Easy Hadoop-Mongo Integration
  - Better Hadoop-Mongo Integration

- Document Oriented
- Organized into Databases and Collections (like Tables)
- JSON-like (BSON)
- Schemaless
- Dynamic, Strong Typing
- Database can “reach into” objects

```
db.people.insert({  
  _id: "mstearn",  
  name: "Mathias Stearn",  
  karma: 42,  
  active: true,  
  birthdate: new Date(517896000000),  
  interests: ["MongoDB", "Python", "Üñíçøďě"],  
  subobject: {foo: "bar"}  
});
```

## JavaScript used for:

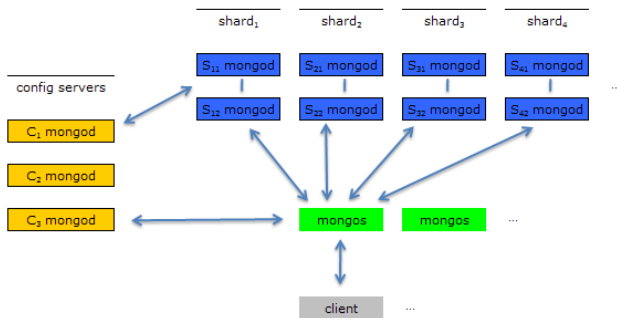
- Shell and Documentation
- (Very) Advanced Queries
- “Group By” Queries
- MapReduce

```
db.users.find({$where: "this.a + this.b >= 42"});  
db.posts.group(  
  { key: "user"  
    , initial: {count:0, comments:0}  
    , reduce: function(doc,out){  
      out.count++;  
      out.comments += doc.comments.length; }  
    , finalize: function(out){  
      out.avg = out.comments / out.count; }  
  });
```

## Fast, Scalable, Available, and Reliable

- Master-Slave replication for Availability and Reliability
  - Replica-Pairs support auto-negotiation for master
- Auto-Sharding for Horizontal Scalability
  - Distributes based on specified field
  - Currently alpha
- MMAP database files to automatically use available RAM
- Asynchronous modifications





## 1 What is MongoDB?

- Document Oriented
- JavaScript Enabled
- Fast, Scalable, Available, and Reliable

## 2 What Makes Mongo Special?

- Native Language Integration
- Rich Data Types
- Atomic Modifiers
- Dynamic Queries

## 3 MapReduce

- Built-In MapReduce
- Easy Hadoop-Mongo Integration
- Better Hadoop-Mongo Integration

## Official

- Java/JVM
- Python
- Ruby
- C/C++
- Perl
- PHP

## Community Supported

Closure, Scala, C#, Haskell, Erlang, *and More*

## JSON

- String (UTF8)
- Double
- Object (hash/map/dict)
- Array
- Bool
- Null / Undefined

## Extras

- Date
- Int32 / Int64
- ObjectID (12 bytes: timestamp + host + pid + counter)
- Binary (with type byte)

- \$set
- \$inc
- \$multiply (soon)
- \$push / \$pushAll
- \$pull / \$pullAll

```
db.posts.update({_id:SOMEID}, {$push:{tags:"mongodb"}})
db.tags.update({_id:"mongodb"}, {$inc:{count:1}},
               {upsert:true})
```

## Simple

```
db.posts.findOne({ user: "mstearn" });  
  
var cursor = db.posts.find({ user: "mstearn" });  
cursor.forEach(function() {  
    doSomething(this.text);  
});
```

## Sorted

```
db.posts.find(  
  { user: "mstearn" }  
) .sort({timestamp:-1})
```

## Paginated

```
db.posts.find(  
  { user: "mstearn" }  
) .sort({timestamp:-1}) .skip(10) .limit(10);
```



## Simple Tag Search

```
db.posts.find(  
  { user: "mstearn"  
    , tags: "mongo"  
  }  
) .sort({timestamp:-1}) .skip(10) .limit(10);
```

## Complex Tag Search

```
db.posts.find(  
  { user: "mstearn"  
    , tags: {$in: ["mongo", "mongodb"]}  
  }  
) .sort({timestamp:-1}) .skip(10) .limit(10);
```

## Nested Objects

```
db.posts.find(  
  { user: "mstearn"  
    , tags: {$in: ["mongo", "mongodb"]}  
    , comments.user: "mdirolf"  
  }  
) .sort({timestamp:-1}) .skip(10) .limit(10);
```

## Regular Expressions

```
db.posts.find(  
  { user: "mstearn"  
    , tags: {$in: ["mongo", "mongodb"]}  
    , comments.user: "mdiroolf"  
    , text: /windows/i  
  }  
) .sort({timestamp:-1}) .skip(10) .limit(10);
```

## Ranges

```
db.posts.find(  
  { user: "mstearn"  
    , tags: {$in: ["mongo", "mongodb"]}  
    , comments.user: "mdirolf"  
    , text: /windows/i  
    , points: {$gt: 10, $lt: 100}  
  }  
) .sort({timestamp:-1}) .skip(10) .limit(10);
```

## Arbitrary JavaScript

```
db.posts.find(  
  { user: "mstearn"  
    , tags: {$in: ["mongo", "mongodb"]}  
    , comments.user: "mdirolf"  
    , text: /windows/i  
    , points: {$gt: 10, $lt 100}  
    , $where: "this.a + this.b >= 42"  
  }  
) .sort({timestamp:-1}) .skip(10) .limit(10);
```

## 1 What is MongoDB?

- Document Oriented
- JavaScript Enabled
- Fast, Scalable, Available, and Reliable

## 2 What Makes Mongo Special?

- Native Language Integration
- Rich Data Types
- Atomic Modifiers
- Dynamic Queries

## 3 MapReduce

- Built-In MapReduce
- Easy Hadoop-Mongo Integration
- Better Hadoop-Mongo Integration

```
db.posts.mapReduce(  
  function() {  
    this.comments.forEach(c) {  
      emit(c.user,  
          {count:1, words:c.text.split().length; } }  
    , function(key, values){  
      for (var i=1; i<values.length; i++){  
        values[0].count += values[i].count;  
        values[0].words += values[i].words; }  
      return values[0]; }  
    , { finalize: function(out){  
        out.avg = out.words / out.count;  
        return out; }  
      , query: {posted: {$gt: new Date(2010,0,1)}}  
      , out: 'posts.comment_stats'  
    }  
  });
```



## Easy Hadoop-Mongo Integration

- mongoexport can export to JSON/CSV/TSV
  - Can also easily use a custom script
- Process in Hadoop
- Use mongoimport to get data back into MongoDB

## Better Hadoop-Mongo Integration

- mongodump writes a stream of BSON to a file
- Write an InputFilter and RecordReader to read BSON
- Write a BSONWriter class to directly use the data
  - Just added two methods to driver to make this easier
- Process the data with the Java/Scala/Closure driver
- Write a custom RecordWriter to either:
  - Dump to a file and use mongorestore
  - Dump the output directly to MongoDB
- Optional: use renameCollection to mimic our MapReduce

## Upcoming events

- NoSQL Live! from Boston (March 11)
- MongoDB Training in San Francisco (March 25)
- San Fransisco MySQL Meetup (April 12)

## Links

- <http://mongo.kylebanker.com> (Try mongo in your browser)
- <http://www.mongodb.org>
- #mongodb on irc.freenode.net
- mongodb-user on google groups
- mathias@10gen.com
- @mathias\_mongo on twitter