

# Document-Oriented DBs and MongoDB

Mathias Stearn

10gen

VolcaNoSQL EU – April 20, 2010

## 1 Document Oriented Databases

- What they are
- What they aren't

## 2 CouchDB

- Pros and Cons

## 3 MongoDB

- Compared to CouchDB
- Intro to Mongo
- Sharding

## 1 Document Oriented Databases

- What they are
- What they aren't

## 2 CouchDB

- Pros and Cons

## 3 MongoDB

- Compared to CouchDB
- Intro to Mongo
- Sharding

- Only Two Document Oriented DBs Right Now
  - MongoDB and CouchDB
  - Everyone has an opinion on what makes a DODB
    - Hard to choose “defining” characteristics
    - This is my take on the space – deal with it

- Document Oriented

- Think JSON Documents, not Word/OOo Documents
- Can store files through Attachments and GridFS
- Could use XML but XML sucks

```
{
  _id: "mstearn",
  name: "Mathias_Stearn",
  karma: 42,
  active: true,
  birthdate: new Date(517896000000),
  interests: ["MongoDB", "Python", "Üñíçøðě"],
  subobjects: [{foo: "bar"},
                {foo: "baz", count: 13}]
}
```

- Hierarchical

- Can nest objects to arbitrary depth
- Server can reach into objects
- Whole “Object” stored at one place on disk

```
{
  comments: [
    { by: 'mstearn', body: 'text', tags: ['empty'],
      votes: {good: 100, bad: 10, net: 90} },
    { by: 'mdirolf', body: 'what?', tags: ['question'],
      votes: {good: 30, bad: 40, net: -10} }
  ]
}
```

- Not Relational
  - Not forced into rows/columns/tables
  - No built-in joins
    - Less need because objects can directly store lists
  - Many-to-Many still possible (learn how at workshop)
  - No SQL (no SQL injections either)
  - No Object-Relational impedance mismatch

- Not Just Key-Value Store
  - Key and value are not separate
  - Supports queries on non-primary keys
    - Secondary Indexes
  - Supports Aggregation
    - Currently via JavaScript MapReduce
    - Both DBs looking into alternatives
  - Can be as fast as a KV store if you only need KV features
    - But still have access to a real database when needed
  - Less custom code needed



- Not the same as stuffing a JSON blob in a database
  - Database understands document format
  - Can query on any field
  - “Use the right tool for the job”

## 1 Document Oriented Databases

- What they are
- What they aren't

## 2 CouchDB

- Pros and Cons

## 3 MongoDB

- Compared to CouchDB
- Into to Mongo
- Sharding

**WARNING**

I am not an expert on CouchDB!

## Pros

- HTTP RESTful Interface
- Stores and communicates in plain JSON
- Query using precomputed JS Map/Reduce views
- Fastest if you use Bulk Insert
- Uses Append-Only File

## Cons

- HTTP RESTful Interface
- Stores and communicates in plain JSON
- Query using precomputed JS Map/Reduce views
- Fastest if you use Bulk Insert
- Uses Append-Only File

## 1 Document Oriented Databases

- What they are
- What they aren't

## 2 CouchDB

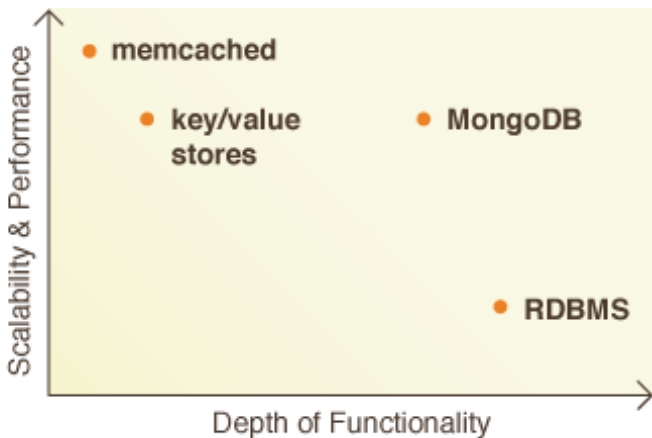
- Pros and Cons

## 3 MongoDB

- Compared to CouchDB
- Into to Mongo
- Sharding

## MongoDB

- Custom wire protocol with many supported languages
- Stores and communicates in BSON (Binary JSON)
- Rich Ad-Hoc Query Language
  - MapReduce for aggregation
- Bulk Insert available, but regular insert is very fast
- Data is updated in place





- The Mongo Shell
  - `http://try.mongodb.org` ← go here now
- Full JS shell + MongoDB extensions
- Most MongoDB documentation uses shell syntax

```

1 db.users.insert({_id: 'mstearn',
2                   name: {first: 'Mathias',
3                           last: 'Stearn'}
4                   company: '10gen',
5                   knows: [ 'MongoDB', 'Python', 'C++' ],
6                   posts: 42})
7
8 db.users.find({_id: 'mstearn'})
9 db.users.find({company: '10gen'})
10 db.users.find({posts: {$gte: 40, $lt: 50}})
11 db.users.find({'name.last': 'Stearn'})
12 db.users.ensureIndex({knows: 1})
13 db.users.find({knows: 'MongoDB'})
14 db.users.find({knows: {$in: [ 'MongoDB', 'Mongo' ]}})
15 db.users.find({knows: {$all: [ 'MongoDB', 'Python' ]}})
16 db.users.find({knows: /^Mongo/})
17 db.users.find().sort({posts: -1}).skip(10).limit(10)

```

## Shameless self-promotion

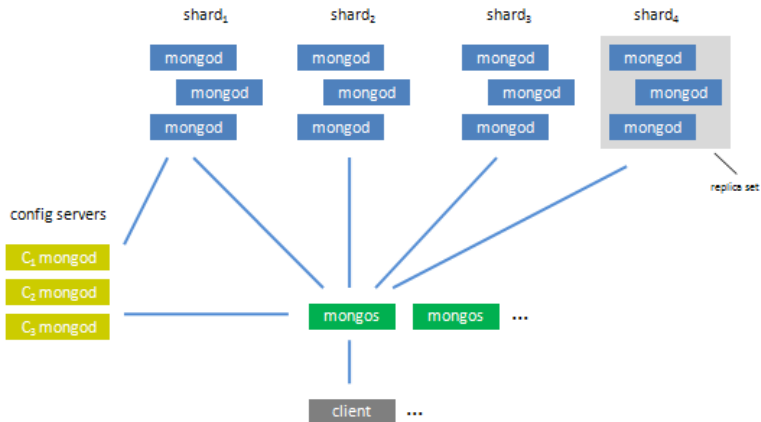
<http://github.com/RedBeard0531/MongoMagic/>

```
1 db.users.find( M._id == 'mstearn' )
2 db.users.find( M.company == '10gen' )
3 db.users.find( 40 <= M.posts < 50 )
4 db.users.find( M.name.last == 'Stearn' )
5 db.users.find( M.knows.IN( 'MongoDB', 'Mongo' ) )
6 db.users.find( M.knows.ALL( 'MongoDB', 'Python' ) )
7 db.users.find( M.knows.STARTSWITH( 'Mongo' ) )
```

```
1 db.zips.insert({_id: '10011', loc: [43, -74]})
2
3 db.zips.ensureIndex({loc: '2d'})
4
5 db.zips.find({loc: {$near: [43, -74]}})
6
7 var box = [[x1, y1], [x2, y2]]
8 db.zips.find({loc: {$within: {$box: box}}})
9
10 var circle = [[x,y], radius]
11 db.zips.find({loc: {$within: {$center: circle}}})
```

```
1 db.posts.insert({_id: ObjectId(123),  
2                 by: 'mstearn',  
3                 title: 'Why_MongoDB_is_Awesome',  
4                 body: 'It_just_is_MASSIVE TYPO',  
5                 tags: [] })  
6  
7 db.posts.update({_id: ObjectId(123)},  
8                 {$set: {body: 'It_just_is' }})  
9  
10 db.posts.update({_id: ObjectId(123)},  
11                 {$push: {tags: 'Citation_Needed' }})  
12  
13 db.tags.update({_id: 'Citation_Needed'},  
14                {$inc: {count: 1}},  
15                {upsert: true})
```

- You (probably) don't need sharding!
- My desktop can handle 100k inserts/s or 160k queries/s
  - For comparison, linkedin gets 55M views/day [quantcast]
  - $55000000 / 8 / 60 / 60 = 1736$  views/sec
- Largest Mongo install is 12TB on single server
  - Mostly large objects
- Wordnik.com has 1.5TB in over 5 Billion docs
  - Sustained 100,000 inserts per second during loading
  - Queries are 4x faster than old MySQL setup
- Speed and Scalability are different things
  - But you only need scalability if you're too slow



# Questions?

## Links

- <http://media.mongodb.org/zijs.json> (for workshop)
- <http://try.mongodb.org> (Try mongo in your browser)
- <http://www.mongodb.org>
- #mongodb on irc.freenode.net
- mongodb-user on google groups

## Contact

- [mathias@10gen.com](mailto:mathias@10gen.com)
- @mathias\_mongo ← follow me!